

Getting and cleaning Data Swirl exercises

Hariharan

This PDF contains the code and output generated for swirl exercises in the course Find the Getting-and-cleaning-data-swirl.Rmd in the same folder as this file to interact with the code and make changes for a better learning experience

1. Manipulating Data with dplyr

Setting up the environment

```
path2csv<-"C:/Users/MAHE/Documents/R/win-library/3.6/swirl/Courses/Getting_and_Cleaning_Data/Manipulating_Data/Getting_and_Cleaning_Data.csv"
mydf <- read.csv(path2csv, stringsAsFactors = FALSE)
dim(mydf)
```

```
## [1] 225468      11
```

```
head(mydf)
```

```
##      X      date      time    size r_version r_arch    r_os    package version
## 1 1 2014-07-08 00:54:41  80589      3.1.0 x86_64  mingw32  htmltools  0.2.4
## 2 2 2014-07-08 00:59:53 321767      3.1.0 x86_64  mingw32    tseries 0.10-32
## 3 3 2014-07-08 00:47:13 748063      3.1.0 x86_64 linux-gnu    party   1.0-15
## 4 4 2014-07-08 00:48:05 606104      3.1.0 x86_64 linux-gnu    Hmisc   3.14-4
## 5 5 2014-07-08 00:46:50  79825      3.0.2 x86_64 linux-gnu    digest   0.6.4
## 6 6 2014-07-08 00:48:04  77681      3.1.0 x86_64 linux-gnu randomForest 4.6-7
##      country ip_id
## 1         US      1
## 2         US      2
## 3         US      3
## 4         US      3
## 5         CA      4
## 6         US      3
```

```
# loading dplyr
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 3.6.3
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
packageVersion("dplyr")
```

```
## [1] '0.8.5'
```

The first step of working with data in dplyr is to load the data into what the package authors call a ‘data frame tbl’ or ‘tbl_df’. Use the following code to create a new tbl_df called cran:

```
cran <- tbl_df(mydf)
rm("mydf")
cran
```

```
## # A tibble: 225,468 x 11
##       X date   time      size r_version r_arch r_os  package version country ip_id
##   <int> <chr> <chr>   <int> <chr>   <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1     1  2014~ 00:5~ 8.06e4 3.1.0   x86_64 ming~ htmlto~ 0.2.4   US        1
## 2     2  2014~ 00:5~ 3.22e5 3.1.0   x86_64 ming~ tseries 0.10-32 US        2
## 3     3  2014~ 00:4~ 7.48e5 3.1.0   x86_64 linu~ party   1.0-15 US        3
## 4     4  2014~ 00:4~ 6.06e5 3.1.0   x86_64 linu~ Hmisc   3.14-4 US        3
## 5     5  2014~ 00:4~ 7.98e4 3.0.2   x86_64 linu~ digest  0.6.4   CA        4
## 6     6  2014~ 00:4~ 7.77e4 3.1.0   x86_64 linu~ random~ 4.6-7   US        3
## 7     7  2014~ 00:4~ 3.94e5 3.1.0   x86_64 linu~ plyr    1.8.1   US        3
## 8     8  2014~ 00:4~ 2.82e4 3.0.2   x86_64 linu~ whisker 0.3-2   US        5
## 9     9  2014~ 00:5~ 5.93e3 <NA>    <NA>    <NA>    Rcpp    0.10.4  CN        6
## 10    10  2014~ 00:1~ 2.21e6 3.0.2   x86_64 linu~ hfligh~ 0.1     US        7
## # ... with 225,458 more rows
```

using select

```
select(cran, ip_id, package, country)
```

```
## # A tibble: 225,468 x 3
##   ip_id package      country
##   <int> <chr>         <chr>
## 1     1  htmltools     US
## 2     2  tseries       US
## 3     3  party         US
## 4     3  Hmisc         US
## 5     4  digest        CA
## 6     3  randomForest  US
## 7     3  plyr          US
## 8     5  whisker       US
## 9     6  Rcpp          CN
## 10    7  hflights      US
## # ... with 225,458 more rows
```

```
select(cran, r_arch:country)
```

```
## # A tibble: 225,468 x 5
##   r_arch r_os      package      version country
##   <chr>  <chr>    <chr>      <chr>    <chr>
## 1 x86_64 mingw32  htmltools  0.2.4    US
## 2 x86_64 mingw32  tseries    0.10-32  US
## 3 x86_64 linux-gnu party        1.0-15   US
## 4 x86_64 linux-gnu Hmisc       3.14-4   US
## 5 x86_64 linux-gnu digest      0.6.4    CA
## 6 x86_64 linux-gnu randomForest 4.6-7    US
## 7 x86_64 linux-gnu plyr         1.8.1    US
## 8 x86_64 linux-gnu whisker      0.3-2    US
## 9 <NA>    <NA>      Rcpp       0.10.4   CN
## 10 x86_64 linux-gnu hflights    0.1      US
## # ... with 225,458 more rows
```

```
select(cran, -time)
```

```
## # A tibble: 225,468 x 10
##       X date      size r_version r_arch r_os  package  version country ip_id
##   <int> <chr>    <int> <chr>    <chr> <chr>  <chr>    <chr>    <chr>  <int>
## 1     1 2014-0~  80589 3.1.0    x86_64 mingw~ htmltools 0.2.4    US        1
## 2     2 2014-0~  321767 3.1.0    x86_64 mingw~ tseries   0.10-32  US        2
## 3     3 2014-0~  748063 3.1.0    x86_64 linux~ party     1.0-15   US        3
## 4     4 2014-0~  606104 3.1.0    x86_64 linux~ Hmisc     3.14-4   US        3
## 5     5 2014-0~  79825 3.0.2    x86_64 linux~ digest    0.6.4    CA        4
## 6     6 2014-0~  77681 3.1.0    x86_64 linux~ randomFo~ 4.6-7    US        3
## 7     7 2014-0~  393754 3.1.0    x86_64 linux~ plyr      1.8.1    US        3
## 8     8 2014-0~  28216 3.0.2    x86_64 linux~ whisker   0.3-2    US        5
## 9     9 2014-0~   5928 <NA>     <NA>    <NA>    Rcpp      0.10.4   CN        6
## 10    10 2014-0~ 2206029 3.0.2    x86_64 linux~ hflights  0.1      US        7
## # ... with 225,458 more rows
```

```
select(cran, -(X:size))
```

```
## # A tibble: 225,468 x 7
##   r_version r_arch r_os      package      version country ip_id
##   <chr>    <chr>  <chr>    <chr>      <chr>    <chr>    <int>
## 1 3.1.0    x86_64 mingw32  htmltools  0.2.4    US        1
## 2 3.1.0    x86_64 mingw32  tseries    0.10-32  US        2
## 3 3.1.0    x86_64 linux-gnu party        1.0-15   US        3
## 4 3.1.0    x86_64 linux-gnu Hmisc       3.14-4   US        3
## 5 3.0.2    x86_64 linux-gnu digest      0.6.4    CA        4
## 6 3.1.0    x86_64 linux-gnu randomForest 4.6-7    US        3
## 7 3.1.0    x86_64 linux-gnu plyr         1.8.1    US        3
## 8 3.0.2    x86_64 linux-gnu whisker      0.3-2    US        5
## 9 <NA>     <NA>    <NA>      Rcpp       0.10.4   CN        6
## 10 3.0.2    x86_64 linux-gnu hflights    0.1      US        7
## # ... with 225,458 more rows
```

using filter

```
filter(cran, package == "swirl")
```

```
## # A tibble: 820 x 11
##       X date   time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>   <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1    27 2014~ 00:1~ 105350 3.0.2   x86_64 ming~ swirl 2.2.9   US        20
## 2   156 2014~ 00:2~  41261 3.1.0   x86_64 linu~ swirl 2.2.9   US         66
## 3   358 2014~ 00:1~ 105335 2.15.2  x86_64 ming~ swirl 2.2.9   CA       115
## 4   593 2014~ 00:5~ 105465 3.1.0   x86_64 darw~ swirl 2.2.9   MX       162
## 5   831 2014~ 00:5~ 105335 3.0.3   x86_64 ming~ swirl 2.2.9   US        57
## 6   997 2014~ 00:3~  41261 3.1.0   x86_64 ming~ swirl 2.2.9   US        70
## 7  1023 2014~ 00:3~ 106393 3.1.0   x86_64 ming~ swirl 2.2.9   BR       248
## 8  1144 2014~ 00:0~ 106534 3.0.2   x86_64 linu~ swirl 2.2.9   US       261
## 9  1402 2014~ 00:4~  41261 3.1.0   i386  ming~ swirl 2.2.9   US       234
## 10 1424 2014~ 00:4~ 106393 3.1.0   x86_64 linu~ swirl 2.2.9   US       301
## # ... with 810 more rows
```

```
filter(cran, r_version == "3.1.1", country == "US")
```

```
## # A tibble: 1,588 x 11
##       X date   time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>   <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1  2216 2014~ 00:4~  3.85e5 3.1.1   x86_64 darw~ colors~ 1.2-4   US       191
## 2 17332 2014~ 03:3~  1.97e5 3.1.1   x86_64 darw~ httr    0.3     US     1704
## 3 17465 2014~ 03:2~  2.33e4 3.1.1   x86_64 darw~ snow    0.3-13 US        62
## 4 18844 2014~ 03:5~  1.91e5 3.1.1   x86_64 darw~ maxLik  1.2-0   US    1533
## 5 30182 2014~ 04:1~  7.77e4 3.1.1   i386  ming~ random~ 4.6-7   US       646
## 6 30193 2014~ 04:0~  2.35e6 3.1.1   i386  ming~ ggplot2 1.0.0   US         8
## 7 30195 2014~ 04:0~  2.99e5 3.1.1   i386  ming~ fExtre~ 3010.81 US     2010
## 8 30217 2014~ 04:3~  5.68e5 3.1.1   i386  ming~ rJava   0.9-6   US        98
## 9 30245 2014~ 04:1~  5.27e5 3.1.1   i386  ming~ LPCM    0.44-8 US         8
## 10 30354 2014~ 04:3~  1.76e6 3.1.1   i386  ming~ mgcv    1.8-1   US    2122
## # ... with 1,578 more rows
```

```
filter(cran, r_version <= "3.0.2", country == "IN")
```

```
## # A tibble: 4,139 x 11
##       X date   time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>   <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1   348 2014~ 00:4~  1.02e7 3.0.0   x86_64 ming~ BH       1.54.0~ IN       112
## 2  9990 2014~ 02:1~  3.97e5 3.0.2   x86_64 linu~ equate~ 1.1     IN     1054
## 3  9991 2014~ 02:1~  1.19e5 3.0.2   x86_64 linu~ ggdend~ 0.1-14  IN     1054
## 4  9992 2014~ 02:1~  8.18e4 3.0.2   x86_64 linu~ dfcrm   0.2-2   IN     1054
## 5 10022 2014~ 02:1~  1.56e6 2.15.0  x86_64 ming~ RcppAr~ 0.4.32~ IN     1060
## 6 10023 2014~ 02:1~  1.18e6 2.15.1  i686  linu~ foreca~ 5.4     IN     1060
## 7 10189 2014~ 02:3~  9.09e5 3.0.2   x86_64 linu~ editru~ 2.7.2   IN     1054
## 8 10199 2014~ 02:3~  1.78e5 3.0.2   x86_64 linu~ energy  1.6.1   IN     1054
## 9 10200 2014~ 02:3~  5.18e4 3.0.2   x86_64 linu~ ENmisc  1.2-7   IN     1054
## 10 10201 2014~ 02:3~  6.52e4 3.0.2   x86_64 linu~ entropy 1.2.0   IN     1054
## # ... with 4,129 more rows
```

```
filter(cran, country == "US" | country == "IN")
```

```
## # A tibble: 95,283 x 11
##       X date time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>    <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1     1 2014~ 00:5~ 8.06e4 3.1.0    x86_64 ming~ htmlto~ 0.2.4   US        1
## 2     2 2014~ 00:5~ 3.22e5 3.1.0    x86_64 ming~ tseries 0.10-32 US        2
## 3     3 2014~ 00:4~ 7.48e5 3.1.0    x86_64 linu~ party   1.0-15 US        3
## 4     4 2014~ 00:4~ 6.06e5 3.1.0    x86_64 linu~ Hmisc   3.14-4  US        3
## 5     6 2014~ 00:4~ 7.77e4 3.1.0    x86_64 linu~ random~ 4.6-7   US        3
## 6     7 2014~ 00:4~ 3.94e5 3.1.0    x86_64 linu~ plyr    1.8.1   US        3
## 7     8 2014~ 00:4~ 2.82e4 3.0.2    x86_64 linu~ whisker 0.3-2   US        5
## 8    10 2014~ 00:1~ 2.21e6 3.0.2    x86_64 linu~ hfligh~ 0.1     US        7
## 9    11 2014~ 00:1~ 5.27e5 3.0.2    x86_64 linu~ LPCM    0.44-8  US        8
## 10   12 2014~ 00:1~ 2.35e6 2.14.1   x86_64 linu~ ggplot2 1.0.0   US        8
## # ... with 95,273 more rows
```

```
filter(cran, size > 100500, r_os == "linux-gnu")
```

```
## # A tibble: 33,683 x 11
##       X date time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>    <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1     3 2014~ 00:4~ 7.48e5 3.1.0    x86_64 linu~ party   1.0-15 US        3
## 2     4 2014~ 00:4~ 6.06e5 3.1.0    x86_64 linu~ Hmisc   3.14-4  US        3
## 3     7 2014~ 00:4~ 3.94e5 3.1.0    x86_64 linu~ plyr    1.8.1   US        3
## 4    10 2014~ 00:1~ 2.21e6 3.0.2    x86_64 linu~ hfligh~ 0.1     US        7
## 5    11 2014~ 00:1~ 5.27e5 3.0.2    x86_64 linu~ LPCM    0.44-8  US        8
## 6    12 2014~ 00:1~ 2.35e6 2.14.1   x86_64 linu~ ggplot2 1.0.0   US        8
## 7    14 2014~ 00:1~ 3.10e6 3.0.2    x86_64 linu~ Rcpp    0.9.7   VE       10
## 8    15 2014~ 00:1~ 5.68e5 3.1.0    x86_64 linu~ rJava   0.9-6   US       11
## 9    16 2014~ 00:1~ 1.60e6 3.1.0    x86_64 linu~ RSQLite 0.11.4  US        7
## 10   18 2014~ 00:2~ 1.87e5 3.1.0    x86_64 linu~ ipred   0.9-3   DE       13
## # ... with 33,673 more rows
```

```
filter(cran, !is.na(r_version))
```

```
## # A tibble: 207,205 x 11
##       X date time      size r_version r_arch r_os package version country ip_id
##   <int> <chr> <chr>   <int> <chr>    <chr> <chr> <chr>   <chr>   <chr>   <int>
## 1     1 2014~ 00:5~ 8.06e4 3.1.0    x86_64 ming~ htmlto~ 0.2.4   US        1
## 2     2 2014~ 00:5~ 3.22e5 3.1.0    x86_64 ming~ tseries 0.10-32 US        2
## 3     3 2014~ 00:4~ 7.48e5 3.1.0    x86_64 linu~ party   1.0-15 US        3
## 4     4 2014~ 00:4~ 6.06e5 3.1.0    x86_64 linu~ Hmisc   3.14-4  US        3
## 5     5 2014~ 00:4~ 7.98e4 3.0.2    x86_64 linu~ digest 0.6.4   CA        4
## 6     6 2014~ 00:4~ 7.77e4 3.1.0    x86_64 linu~ random~ 4.6-7   US        3
## 7     7 2014~ 00:4~ 3.94e5 3.1.0    x86_64 linu~ plyr    1.8.1   US        3
## 8     8 2014~ 00:4~ 2.82e4 3.0.2    x86_64 linu~ whisker 0.3-2   US        5
## 9    10 2014~ 00:1~ 2.21e6 3.0.2    x86_64 linu~ hfligh~ 0.1     US        7
## 10   11 2014~ 00:1~ 5.27e5 3.0.2    x86_64 linu~ LPCM    0.44-8  US        8
## # ... with 207,195 more rows
```

using arrange

```
cran2 <- select(cran, size:ip_id)
arrange(cran2, ip_id)#ip_id is in ascending order
```

```
## # A tibble: 225,468 x 8
##   size r_version r_arch r_os package version country ip_id
##   <int> <chr>    <chr> <chr> <chr>    <chr> <chr>    <int>
## 1  80589 3.1.0    x86_64 mingw32 htmltools 0.2.4 US        1
## 2 180562 3.0.2    x86_64 mingw32 yaml      2.1.13 US        1
## 3 190120 3.1.0    i386    mingw32 babel     0.2-6 US        1
## 4 321767 3.1.0    x86_64 mingw32 tseries   0.10-32 US        2
## 5  52281 3.0.3    x86_64 darwin10.8.0 quadprog 1.5-5 US        2
## 6 876702 3.1.0    x86_64 linux-gnu zoo        1.7-11 US        2
## 7 321764 3.0.2    x86_64 linux-gnu tseries   0.10-32 US        2
## 8 876702 3.1.0    x86_64 linux-gnu zoo        1.7-11 US        2
## 9 321768 3.1.0    x86_64 mingw32 tseries   0.10-32 US        2
## 10 784093 3.1.0    x86_64 linux-gnu strucchange 1.5-0 US        2
## # ... with 225,458 more rows
```

```
arrange(cran2, desc(ip_id))#To do the same, but in descending order
```

```
## # A tibble: 225,468 x 8
##   size r_version r_arch r_os package version country ip_id
##   <int> <chr>    <chr> <chr> <chr>    <chr> <chr>    <int>
## 1   5933 <NA>    <NA> <NA> CPE      1.4.2 CN      13859
## 2 569241 3.1.0    x86_64 mingw32 multcompView 0.1-5 US      13858
## 3 228444 3.1.0    x86_64 mingw32 tourr     0.5.3 NZ      13857
## 4 308962 3.1.0    x86_64 darwin13.1.0 ctv       0.7-9 CN      13856
## 5 950964 3.0.3    i386    mingw32 knitr     1.6 CA      13855
## 6  80185 3.0.3    i386    mingw32 htmltools 0.2.4 CA      13855
## 7 1431750 3.0.3    i386    mingw32 shiny     0.10.0 CA      13855
## 8 2189695 3.1.0    x86_64 mingw32 RMySQL    0.9-3 US      13854
## 9 4818024 3.1.0    i386    mingw32 igraph    0.7.1 US      13853
## 10 197495 3.1.0    x86_64 mingw32 coda      0.16-1 US      13852
## # ... with 225,458 more rows
```

```
arrange(cran2, package, ip_id)# first arrange by package names (ascending alphabetically), then by ip_id
```

```
## # A tibble: 225,468 x 8
##   size r_version r_arch r_os package version country ip_id
##   <int> <chr>    <chr> <chr> <chr>    <chr> <chr>    <int>
## 1 71677 3.0.3    x86_64 darwin10.8.0 A3      0.9.2 CN      1003
## 2 71672 3.1.0    x86_64 linux-gnu A3      0.9.2 US      1015
## 3 71677 3.1.0    x86_64 mingw32 A3      0.9.2 IN      1054
## 4 70438 3.0.1    x86_64 darwin10.8.0 A3      0.9.2 CN      1513
## 5 71677 <NA>    <NA> <NA> A3      0.9.2 BR      1526
## 6 71892 3.0.2    x86_64 linux-gnu A3      0.9.2 IN      1542
## 7 71677 3.1.0    x86_64 linux-gnu A3      0.9.2 ZA      2925
## 8 71672 3.1.0    x86_64 mingw32 A3      0.9.2 IL      3889
## 9 71677 3.0.3    x86_64 mingw32 A3      0.9.2 DE      3917
```

```
## 10 71672 3.1.0      x86_64 mingw32      A3      0.9.2  US      4219
## # ... with 225,458 more rows
```

```
arrange(cran2, country, desc(r_version), ip_id)
```

```
## # A tibble: 225,468 x 8
##       size r_version r_arch r_os      package      version  country ip_id
##   <int> <chr>      <chr> <chr>    <chr>      <chr>    <chr>  <int>
## 1 1556858 3.1.1      i386  mingw32 RcppArmadillo 0.4.320.0 A1      2843
## 2 1823512 3.1.0      x86_64 linux-gnu mgcv          1.8-1      A1      2843
## 3 15732 3.1.0      i686  linux-gnu grnn          0.1.0      A1      3146
## 4 3014840 3.1.0      x86_64 mingw32 Rcpp          0.11.2     A1      3146
## 5 660087 3.1.0      i386  mingw32 xts           0.9-7      A1      3146
## 6 522261 3.1.0      i386  mingw32 FNN           1.1        A1      3146
## 7 522263 3.1.0      i386  mingw32 FNN           1.1        A1      3146
## 8 1676627 3.1.0      x86_64 linux-gnu rgeos         0.3-5      A1      3146
## 9 2118530 3.1.0      x86_64 linux-gnu spacetime     1.1-0      A1      3146
## 10 2217180 3.1.0      x86_64 mingw32 gstat         1.0-19     A1      3146
## # ... with 225,458 more rows
```

using mutate

```
cran3 <- select(cran, ip_id, package, size)
mutate(cran3, size_mb = size / 2^20)
```

```
## # A tibble: 225,468 x 4
##       ip_id package      size size_mb
##   <int> <chr>      <int>  <dbl>
## 1     1  htmltools    80589 0.0769
## 2     2  tseries    321767 0.307
## 3     3  party      748063 0.713
## 4     3  Hmisc      606104 0.578
## 5     4  digest      79825 0.0761
## 6     3  randomForest 77681 0.0741
## 7     3  plyr       393754 0.376
## 8     5  whisker     28216 0.0269
## 9     6  Rcpp        5928 0.00565
## 10    7  hflights   2206029 2.10
## # ... with 225,458 more rows
```

```
mutate(cran3, size_mb = size / 2^20, size_gb = size_mb / 2^10)
```

```
## # A tibble: 225,468 x 5
##       ip_id package      size size_mb  size_gb
##   <int> <chr>      <int>  <dbl>  <dbl>
## 1     1  htmltools    80589 0.0769 0.0000751
## 2     2  tseries    321767 0.307 0.000300
## 3     3  party      748063 0.713 0.000697
## 4     3  Hmisc      606104 0.578 0.000564
## 5     4  digest      79825 0.0761 0.0000743
```

```
## 6      3 randomForest    77681 0.0741 0.0000723
## 7      3 plyr           393754 0.376 0.000367
## 8      5 whisker        28216 0.0269 0.0000263
## 9      6 Rcpp           5928 0.00565 0.00000552
## 10     7 hflights       2206029 2.10 0.00205
## # ... with 225,458 more rows
```

```
mutate(cran3, correct_size = size + 1000)
```

```
## # A tibble: 225,468 x 4
##   ip_id package      size correct_size
##   <int> <chr>      <int>      <dbl>
## 1     1 1 htmltools    80589      81589
## 2     2 2 tseries     321767     322767
## 3     3 3 party       748063     749063
## 4     4 3 Hmisc       606104     607104
## 5     5 4 digest       79825      80825
## 6     6 3 randomForest 77681      78681
## 7     7 3 plyr       393754     394754
## 8     8 5 whisker     28216      29216
## 9     9 6 Rcpp        5928       6928
## 10    10 7 hflights    2206029    2207029
## # ... with 225,458 more rows
```

using summarize

```
summarize(cran, avg_bytes = mean(size))
```

```
## # A tibble: 1 x 1
##   avg_bytes
##   <dbl>
## 1 844086.
```

2. Grouping and Chaining with dplyr

Setting up the environment

```
path2csv<-"C:/Users/MAHE/Documents/R/win-library/3.6/swirl/Courses/Getting_and_Cleaning_Data/Manipulation"
mydf <- read.csv(path2csv, stringsAsFactors = FALSE)
dim(mydf)
```

```
## [1] 225468      11
```

```
head(mydf)
```



```
##      X      date      time      size r_version r_arch      r_os      package version
## 1 1 2014-07-08 00:54:41 80589      3.1.0 x86_64 mingw32      htmltools 0.2.4
## 2 2 2014-07-08 00:59:53 321767      3.1.0 x86_64 mingw32      tseries 0.10-32
## 3 3 2014-07-08 00:47:13 748063      3.1.0 x86_64 linux-gnu      party 1.0-15
## 4 4 2014-07-08 00:48:05 606104      3.1.0 x86_64 linux-gnu      Hmisc 3.14-4
## 5 5 2014-07-08 00:46:50 79825      3.0.2 x86_64 linux-gnu      digest 0.6.4
## 6 6 2014-07-08 00:48:04 77681      3.1.0 x86_64 linux-gnu randomForest 4.6-7
##      country ip_id
## 1      US      1
## 2      US      2
## 3      US      3
## 4      US      3
## 5      CA      4
## 6      US      3
```

```
# loading dplyr
library(dplyr)
packageVersion("dplyr")
```

```
## [1] '0.8.5'
```

```
cran <- tbl_df(mydf)
```

using group-by

```
by_package <- group_by(cran, package)
by_package
```

```
## # A tibble: 225,468 x 11
## # Groups:   package [6,023]
##      X date      time      size r_version r_arch r_os      package version country ip_id
##      <int> <chr> <chr> <int> <chr> <chr> <chr> <chr> <chr> <chr> <int>
## 1      1 2014~ 00:5~ 8.06e4 3.1.0      x86_64 ming~ htmlto~ 0.2.4   US      1
## 2      2 2014~ 00:5~ 3.22e5 3.1.0      x86_64 ming~ tseries 0.10-32 US      2
## 3      3 2014~ 00:4~ 7.48e5 3.1.0      x86_64 linu~ party   1.0-15 US      3
## 4      4 2014~ 00:4~ 6.06e5 3.1.0      x86_64 linu~ Hmisc   3.14-4 US      3
## 5      5 2014~ 00:4~ 7.98e4 3.0.2      x86_64 linu~ digest  0.6.4   CA      4
## 6      6 2014~ 00:4~ 7.77e4 3.1.0      x86_64 linu~ random~ 4.6-7   US      3
## 7      7 2014~ 00:4~ 3.94e5 3.1.0      x86_64 linu~ plyr    1.8.1   US      3
## 8      8 2014~ 00:4~ 2.82e4 3.0.2      x86_64 linu~ whisker 0.3-2   US      5
## 9      9 2014~ 00:5~ 5.93e3 <NA>      <NA>    <NA>    Rcpp    0.10.4  CN      6
## 10     10 2014~ 00:1~ 2.21e6 3.0.2      x86_64 linu~ hfligh~ 0.1     US      7
## # ... with 225,458 more rows
```

```
summarize(by_package, mean(size))
```

```
## # A tibble: 6,023 x 2
##      package      `mean(size)`
##      <chr>          <dbl>
```

```
## 1 A3 62195.
## 2 abc 4826665
## 3 abcdeFBA 455980.
## 4 ABCExtremes 22904.
## 5 ABCoptim 17807.
## 6 ABCp2 30473.
## 7 abctools 2589394
## 8 abd 453631.
## 9 abf2 35693.
## 10 abind 32939.
## # ... with 6,013 more rows
```

```
pack_sum <- summarize(by_package,
                      count = n(),
                      unique = n_distinct(ip_id),
                      countries = n_distinct(country),
                      avg_bytes = mean(size))
pack_sum
```

```
## # A tibble: 6,023 x 5
##   package      count unique countries avg_bytes
##   <chr>      <int>  <int>    <int>    <dbl>
## 1 A3         25     24      10    62195.
## 2 abc        29     25      16  4826665
## 3 abcdeFBA    15     15       9   455980.
## 4 ABCExtremes 18     17       9   22904.
## 5 ABCoptim    16     15       9   17807.
## 6 ABCp2       18     17      10   30473.
## 7 abctools    19     19      11  2589394
## 8 abd        17     16      10   453631.
## 9 abf2       13     13       9   35693.
## 10 abind     396    365      50   32939.
## # ... with 6,013 more rows
```

The ‘count’ column, created with `n()`, contains the total number of rows (i.e. downloads) for each package. The ‘unique’ column, created with `n_distinct(ip_id)`, gives the total number of unique downloads for each package, as measured by the number of distinct `ip_id`’s. The ‘countries’ column, created with `n_distinct(country)`, provides the number of countries in which each package was downloaded. And finally, the ‘avg_bytes’ column, created with `mean(size)`, contains the mean download size (in bytes) for each package.

Naturally, we’d like to know which packages were most popular on the day these data were collected (July 8, 2014). Let’s start by isolating the top 1% of packages, based on the total number of downloads as measured by the ‘count’ column.

```
quantile(pack_sum$count, probs = 0.99)
```

```
## 99%
## 679.56
```

```
top_counts <- filter(pack_sum, count > 679)
top_counts
```

```
## # A tibble: 61 x 5
##   package    count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 bitops      1549  1408         76   28715.
## 2 car          1008   837         64  1229122.
## 3 caTools       812   699         64   176589.
## 4 colorspace  1683  1433         80   357411.
## 5 data.table   680   564         59  1252721.
## 6 DBI          2599   492         48  206933.
## 7 devtools     769   560         55   212933.
## 8 dichromat   1486  1257         74   134732.
## 9 digest       2210  1894         83   120549.
## 10 doSNOW       740    75         24    8364.
## # ... with 51 more rows
```

```
View(top_counts)
top_counts_sorted <- arrange(top_counts, desc(count))
top_counts_sorted
```

```
## # A tibble: 61 x 5
##   package    count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 ggplot2    4602  1680         81  2427716.
## 2 Rcpp       3195  2044         84  2512100.
## 3 plyr       2908  1754         81   799123.
## 4 rJava      2773   963         70   633522.
## 5 DBI        2599   492         48  206933.
## 6 LPCM       2335    17         10   526814.
## 7 stringr    2267  1948         82   65277.
## 8 digest     2210  1894         83   120549.
## 9 reshape2   2032  1652         76   330128.
## 10 foreach   1984   485         53   358070.
## # ... with 51 more rows
```

```
quantile(pack_sum$unique, probs = 0.99)
```

```
## 99%
## 465
```

```
top_unique <- filter(pack_sum, unique > 465)
top_unique_sorted <- arrange(top_unique, desc(unique))
top_unique_sorted
```

```
## # A tibble: 60 x 5
##   package    count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 Rcpp       3195  2044         84  2512100.
## 2 stringr    2267  1948         82   65277.
## 3 digest     2210  1894         83   120549.
## 4 plyr       2908  1754         81   799123.
## 5 ggplot2    4602  1680         81  2427716.
## 6 reshape2   2032  1652         76   330128.
```

```
## 7 RColorBrewer 1890 1584 79 22764.
## 8 colorspace 1683 1433 80 357411.
## 9 bitops 1549 1408 76 28715.
## 10 scales 1726 1408 77 126819.
## # ... with 50 more rows
```

Chaining

```
by_package <- group_by(cran, package)
pack_sum <- summarize(by_package,
  count = n(),
  unique = n_distinct(ip_id),
  countries = n_distinct(country),
  avg_bytes = mean(size))
top_countries <- filter(pack_sum, countries > 60)
result1 <- arrange(top_countries, desc(countries), avg_bytes)
print(result1)
```

```
## # A tibble: 46 x 5
##   package      count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 Rcpp        3195  2044         84  2512100.
## 2 digest      2210  1894         83  120549.
## 3 stringr     2267  1948         82   65277.
## 4 plyr        2908  1754         81   799123.
## 5 ggplot2     4602  1680         81  2427716.
## 6 colorspace  1683  1433         80  357411.
## 7 RColorBrewer 1890  1584         79   22764.
## 8 scales      1726  1408         77  126819.
## 9 bitops      1549  1408         76   28715.
## 10 reshape2   2032  1652         76   330128.
## # ... with 36 more rows
```

Same operations as above but using function call embedding

```
result2 <-
  arrange(
    filter(
      summarize(
        group_by(cran,
          package
        ),
        count = n(),
        unique = n_distinct(ip_id),
        countries = n_distinct(country),
        avg_bytes = mean(size)
      ),
      countries > 60
    ),
    desc(countries),
```

```

    avg_bytes
  )

print(result2)

## # A tibble: 46 x 5
##   package      count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 Rcpp        3195  2044         84  2512100.
## 2 digest      2210  1894         83   120549.
## 3 stringr     2267  1948         82    65277.
## 4 plyr        2908  1754         81   799123.
## 5 ggplot2     4602  1680         81  2427716.
## 6 colorspace  1683  1433         80   357411.
## 7 RColorBrewer 1890  1584         79    22764.
## 8 scales      1726  1408         77   126819.
## 9 bitops      1549  1408         76    28715.
## 10 reshape2   2032  1652         76   330128.
## # ... with 36 more rows

```

In this script, we've used a special chaining operator, %>%

*# you read it, you can pronounce the %>% operator as
the word 'then'.*

```

result3 <-
  cran %>%
  group_by(package) %>%
  summarize(count = n(),
            unique = n_distinct(ip_id),
            countries = n_distinct(country),
            avg_bytes = mean(size)
  ) %>%
  filter(countries > 60) %>%
  arrange(desc(countries), avg_bytes)

# Print result to console
print(result3)

```

```

## # A tibble: 46 x 5
##   package      count unique countries avg_bytes
##   <chr>      <int> <int>      <int>      <dbl>
## 1 Rcpp        3195  2044         84  2512100.
## 2 digest      2210  1894         83   120549.
## 3 stringr     2267  1948         82    65277.
## 4 plyr        2908  1754         81   799123.
## 5 ggplot2     4602  1680         81  2427716.
## 6 colorspace  1683  1433         80   357411.
## 7 RColorBrewer 1890  1584         79    22764.
## 8 scales      1726  1408         77   126819.
## 9 bitops      1549  1408         76    28715.
## 10 reshape2   2032  1652         76   330128.
## # ... with 36 more rows

```

```
# select() the following columns from cran. Keep in mind
# that when you're using the chaining operator, you don't
# need to specify the name of the data tbl in your call to
# select().
#
# 1. ip_id
# 2. country
# 3. package
# 4. size
#
# The call to print() at the end of the chain is optional,
# but necessary if you want your results printed to the
# console. Note that since there are no additional arguments
# to print(), you can leave off the parentheses after
# the function name. This is a convenient feature of the %>%
# operator.
```

```
cran %>%
  select(ip_id, country, package, size) %>%
  print
```

```
## # A tibble: 225,468 x 4
##   ip_id country package      size
##   <int> <chr>   <chr>      <int>
## 1     1  US    htmltools    80589
## 2     2  US    tseries    321767
## 3     3  US    party      748063
## 4     3  US    Hmisc      606104
## 5     4  CA    digest      79825
## 6     3  US    randomForest 77681
## 7     3  US    plyr      393754
## 8     5  US    whisker     28216
## 9     6  CN    Rcpp        5928
## 10    7  US    hflights   2206029
## # ... with 225,458 more rows
```

```
# Use mutate() to add a column called size_mb that contains
# the size of each download in megabytes (i.e. size / 2^20).
#
# If you want your results printed to the console, add
# print to the end of your chain.
```

```
cran %>%
  select(ip_id, country, package, size) %>%
  mutate(size_mb = size / 2^20)
```

```
## # A tibble: 225,468 x 5
##   ip_id country package      size size_mb
##   <int> <chr>   <chr>      <int>   <dbl>
## 1     1  US    htmltools    80589 0.0769
## 2     2  US    tseries    321767 0.307
## 3     3  US    party      748063 0.713
```

```
## 4      3 US      Hmisc          606104 0.578
## 5      4 CA      digest          79825 0.0761
## 6      3 US      randomForest    77681 0.0741
## 7      3 US      plyr           393754 0.376
## 8      5 US      whisker         28216 0.0269
## 9      6 CN      Rcpp            5928 0.00565
## 10     7 US      hflights        2206029 2.10
## # ... with 225,458 more rows
```

```
# Use filter() to select all rows for which size_mb is
# less than or equal to (<=) 0.5.
#
# If you want your results printed to the console, add
# print to the end of your chain.
```

```
cran %>%
  select(ip_id, country, package, size) %>%
  mutate(size_mb = size / 2^20) %>%
  filter(size_mb <= 0.5)
```

```
## # A tibble: 142,021 x 5
##   ip_id country package      size size_mb
##   <int> <chr>   <chr>      <int>  <dbl>
## 1      1 US      htmltools    80589 0.0769
## 2      2 US      tseries    321767 0.307
## 3      4 CA      digest      79825 0.0761
## 4      3 US      randomForest 77681 0.0741
## 5      3 US      plyr       393754 0.376
## 6      5 US      whisker     28216 0.0269
## 7      6 CN      Rcpp        5928 0.00565
## 8     13 DE      ipred      186685 0.178
## 9     14 US      mnormt     36204 0.0345
## 10    16 US      iterators  289972 0.277
## # ... with 142,011 more rows
```

```
# arrange() the result by size_mb, in descending order.
#
# If you want your results printed to the console, add
# print to the end of your chain.
```

```
cran %>%
  select(ip_id, country, package, size) %>%
  mutate(size_mb = size / 2^20) %>%
  filter(size_mb <= 0.5) %>%
  arrange(desc(size_mb))
```

```
## # A tibble: 142,021 x 5
##   ip_id country package      size size_mb
##   <int> <chr>   <chr>      <int>  <dbl>
## 1 11034 DE      phia      524232 0.500
## 2  9643 US      tis       524152 0.500
## 3  1542 IN      RcppSMC    524060 0.500
## 4 12354 US      lessR      523916 0.500
```

```
## 5 12072 US      colorspace      523880 0.500
## 6 2514 KR      depmixS4          523863 0.500
## 7 1111 US      depmixS4          523858 0.500
## 8 8865 CR      depmixS4          523858 0.500
## 9 5908 CN      RcmdrPlugin.KMggplot2 523852 0.500
## 10 12354 US    RcmdrPlugin.KMggplot2 523852 0.500
## # ... with 142,011 more rows
```

3. Tidying Data with tidyr

Setting up

```
library(tidyr)
```

```
## Warning: package 'tidyr' was built under R version 3.6.3
```

```
# recreating dataset used in lesson
```

```
students<- data.frame("grade" =c('A','B','C','D','E') , "male" = as.integer(c(5,4,8,4,5)), "female" = as.integer(c(3,1,6,5,5)))
```

```
students
```

```
##   grade male female
## 1     A     5      3
## 2     B     4      1
## 3     C     8      6
## 4     D     4      5
## 5     E     5      5
```

```
gather(students, sex, count, -grade)
```

```
##   grade    sex count
## 1     A  male     5
## 2     B  male     4
## 3     C  male     8
## 4     D  male     4
## 5     E  male     5
## 6     A female     3
## 7     B female     1
## 8     C female     6
## 9     D female     5
## 10    E female     5
```

```
students2<- data.frame("grade" =c('A','B','C','D','E') , "male_1" = as.integer(c(7,4,7,8,8)), "female_1" = as.integer(c(0,0,6,5,5)))
```

```
students2
```

```
##   grade male_1 female_1 male_2 female_2
## 1     A       7         0       5         8
## 2     B       4         0       5         8
```



```
## 3      C      7      4      5      6
## 4      D      8      2      8      1
## 5      E      8      4      1      0
```

```
res <- gather(students2, sex_class, count, -grade)
res
```

```
##      grade sex_class count
## 1      A   male_1     7
## 2      B   male_1     4
## 3      C   male_1     7
## 4      D   male_1     8
## 5      E   male_1     8
## 6      A  female_1     0
## 7      B  female_1     0
## 8      C  female_1     4
## 9      D  female_1     2
## 10     E  female_1     4
## 11     A   male_2     5
## 12     B   male_2     5
## 13     C   male_2     5
## 14     D   male_2     8
## 15     E   male_2     1
## 16     A  female_2     8
## 17     B  female_2     8
## 18     C  female_2     6
## 19     D  female_2     1
## 20     E  female_2     0
```

```
separate(res, sex_class, c("sex", "class"))
```

```
##      grade  sex class count
## 1      A   male    1     7
## 2      B   male    1     4
## 3      C   male    1     7
## 4      D   male    1     8
## 5      E   male    1     8
## 6      A female    1     0
## 7      B female    1     0
## 8      C female    1     4
## 9      D female    1     2
## 10     E female    1     4
## 11     A   male    2     5
## 12     B   male    2     5
## 13     C   male    2     5
## 14     D   male    2     8
## 15     E   male    2     1
## 16     A female    2     8
## 17     B female    2     8
## 18     C female    2     6
## 19     D female    2     1
## 20     E female    2     0
```

using chaining

```
students2 %>%  
  gather(sex_class, count, -grade) %>%  
  separate(sex_class, c("sex", "class")) %>%  
  print
```

```
##   grade  sex class count  
## 1     A  male     1     7  
## 2     B  male     1     4  
## 3     C  male     1     7  
## 4     D  male     1     8  
## 5     E  male     1     8  
## 6     A female     1     0  
## 7     B female     1     0  
## 8     C female     1     4  
## 9     D female     1     2  
## 10    E female     1     4  
## 11    A  male     2     5  
## 12    B  male     2     5  
## 13    C  male     2     5  
## 14    D  male     2     8  
## 15    E  male     2     1  
## 16    A female     2     8  
## 17    B female     2     8  
## 18    C female     2     6  
## 19    D female     2     1  
## 20    E female     2     0
```

```
students3<- data.frame(  
  "name" = c("Sally","Sally","Jeff","Jeff","Roger","Roger","Karen","Karen","Brian","Brian"),  
  "test" = c("midterm","final","midterm","final","midterm","final","midterm","final","midterm","final"),  
  "class1" = c("A","C",NA,NA,NA,NA,NA,NA,"B","B"),  
  "class2" = c(NA,NA,"D","E","C","A",NA,NA,NA,NA),  
  "class3" = c("B","C",NA,NA,NA,NA,"C","C",NA,NA),  
  "class5" = c(NA,NA,NA,NA,"B","A",NA,NA,"A","C"),  
  stringsAsFactors = FALSE  
)  
students3
```

```
##   name  test class1 class2 class3 class5  
## 1 Sally midterm     A  <NA>      B  <NA>  
## 2 Sally  final     C  <NA>      C  <NA>  
## 3 Jeff midterm  <NA>      D  <NA>  <NA>  
## 4 Jeff  final  <NA>      E  <NA>  <NA>  
## 5 Roger midterm  <NA>      C  <NA>      B  
## 6 Roger  final  <NA>      A  <NA>      A  
## 7 Karen midterm  <NA>  <NA>      C  <NA>  
## 8 Karen  final  <NA>  <NA>      C  <NA>  
## 9 Brian midterm     B  <NA>  <NA>      A  
## 10 Brian  final     B  <NA>  <NA>      C
```

```

# Call gather() to gather the columns class1
# through class5 into a new variable called class.
# The 'key' should be class, and the 'value'
# should be grade.
#
# tidyr makes it easy to reference multiple adjacent
# columns with class1:class5, just like with sequences
# of numbers.
#
# Since each student is only enrolled in two of
# the five possible classes, there are lots of missing
# values (i.e. NAs). Use the argument na.rm = TRUE
# to omit these values from the final result.
#
# Remember that when you're using the %>% operator,
# the value to the left of it gets inserted as the
# first argument to the function on the right.
#
# Consult ?gather and/or ?chain if you get stuck.
#
students3 %>%
  gather(class, grade, class1:class5, na.rm = TRUE) %>%
  print

```

```

##      name      test  class grade
## 1  Sally midterm class1      A
## 2  Sally   final class1      C
## 9  Brian midterm class1      B
## 10 Brian   final class1      B
## 13 Jeff  midterm class2      D
## 14 Jeff   final class2      E
## 15 Roger midterm class2      C
## 16 Roger   final class2      A
## 21 Sally midterm class3      B
## 22 Sally   final class3      C
## 27 Karen midterm class3      C
## 28 Karen   final class3      C
## 35 Roger midterm class5      B
## 36 Roger   final class5      A
## 39 Brian midterm class5      A
## 40 Brian   final class5      C

```

```

# This script builds on the previous one by appending
# a call to spread(), which will allow us to turn the
# values of the test column, midterm and final, into
# column headers (i.e. variables).
#
# You only need to specify two arguments to spread().
# Can you figure out what they are? (Hint: You don't
# have to specify the data argument since we're using
# the %>% operator.
#
students3 %>%

```

```
gather(class, grade, class1:class5, na.rm = TRUE) %>%
spread(test, grade) %>%
print
```

```
##   name  class final midterm
## 1 Brian class1     B       B
## 2 Brian class5     C       A
## 3 Jeff  class2     E       D
## 4 Karen class3     C       C
## 5 Roger class2     A       C
## 6 Roger class5     A       B
## 7 Sally class1     C       A
## 8 Sally class3     C       B
```

```
library(readr)
```

```
## Warning: package 'readr' was built under R version 3.6.3
```

```
parse_number("class5")
```

```
## [1] 5
```

```
# We want the values in the class columns to be
# 1, 2, ..., 5 and not class1, class2, ..., class5.
#
# Use the mutate() function from dplyr along with
# parse_number(). Hint: You can "overwrite" a column
# with mutate() by assigning a new value to the existing
# column instead of creating a new column.
#
# Check out ?mutate and/or ?parse_number if you need
# a refresher.
#
students3 %>%
  gather(class, grade, class1:class5, na.rm = TRUE) %>%
  spread(test, grade) %>%
  mutate(class = parse_number(class)) %>%
  print
```

```
##   name class final midterm
## 1 Brian     1     B       B
## 2 Brian     5     C       A
## 3 Jeff      2     E       D
## 4 Karen     3     C       C
## 5 Roger     2     A       C
## 6 Roger     5     A       B
## 7 Sally     1     C       A
## 8 Sally     3     C       B
```

```
students4 <- data.frame(
  "id" = as.integer(c(168,168,588,588,710,710,731,731,908,908)),
  "name" = c("Brian","Brian","Sally","Sally","Jeff","Jeff","Roger","Roger","Karen","Karen"),
  "sex" = c("F","F","M","M","M","M","F","F","M","M"),
  "class" = as.integer(c(1,5,1,3,2,4,2,5,3,4)),
  "midterm" = c("B","A","A","B","D","A","C","B","C","A"),
  "final" = c("B","C","C","C","E","C","A","A","C","A"),
  stringsAsFactors = FALSE
)
students4
```

```
##      id  name sex class midterm final
## 1  168 Brian  F     1         B     B
## 2  168 Brian  F     5         A     C
## 3  588 Sally  M     1         A     C
## 4  588 Sally  M     3         B     C
## 5  710 Jeff   M     2         D     E
## 6  710 Jeff   M     4         A     C
## 7  731 Roger  F     2         C     A
## 8  731 Roger  F     5         B     A
## 9  908 Karen  M     3         C     C
## 10 908 Karen  M     4         A     A
```

```
# selecting the id, name, and sex column from students4
# and storing the result in student_info.
#
student_info <- students4 %>%
  select(id, name, sex) %>%
  print
```

```
##      id  name sex
## 1  168 Brian  F
## 2  168 Brian  F
## 3  588 Sally  M
## 4  588 Sally  M
## 5  710 Jeff   M
## 6  710 Jeff   M
## 7  731 Roger  F
## 8  731 Roger  F
## 9  908 Karen  M
## 10 908 Karen  M
```

```
# Add a call to unique() below, which will remove
# duplicate rows from student_info.
#
# Like with the call to the print() function below,
# you can omit the parentheses after the function name.
# This is a nice feature of %>% that applies when
# there are no additional arguments to specify.
#
student_info <- students4 %>%
```

```
select(id, name, sex) %>%
  unique %>%
  print
```

```
##      id  name sex
## 1 168 Brian   F
## 3 588 Sally   M
## 5 710  Jeff   M
## 7 731 Roger   F
## 9 908 Karen   M
```

```
# select() the id, class, midterm, and final columns
# (in that order) and store the result in gradebook.
#
gradebook <- students4 %>%
  select(id, class, midterm, final) %>%
  print
```

```
##      id class midterm final
## 1  168     1        B      B
## 2  168     5        A      C
## 3  588     1        A      C
## 4  588     3        B      C
## 5  710     2        D      E
## 6  710     4        A      C
## 7  731     2        C      A
## 8  731     5        B      A
## 9  908     3        C      C
## 10 908     4        A      A
```

```
passed<- data.frame(
  "name" = c("Brian","Roger","Roger","Karen"),
  "class" = as.integer(c(1,2,5,4)),
  "final" = c("B","A","A","A"),
  stringsAsFactors = FALSE
)
failed <- data.frame(
  "name" = c("Brian","Sally","Sally","Jeff","Jeff","Karen"),
  "class" = as.integer(c(5,1,3,2,4,3)),
  "final" = c("C","C","C","E","C","C"),
  stringsAsFactors = FALSE
)
passed
```

```
##      name class final
## 1 Brian     1      B
## 2 Roger     2      A
## 3 Roger     5      A
## 4 Karen     4      A
```

```
failed
```

```
##      name class final
## 1 Brian      5      C
## 2 Sally      1      C
## 3 Sally      3      C
## 4  Jeff      2      E
## 5  Jeff      4      C
## 6 Karen      3      C
```

```
passed <- passed %>% mutate(status = "passed")
failed <- failed %>% mutate(status = "failed")
bind_rows(passed, failed)
```

```
##      name class final status
## 1  Brian      1      B passed
## 2  Roger      2      A passed
## 3  Roger      5      A passed
## 4  Karen      4      A passed
## 5  Brian      5      C failed
## 6  Sally      1      C failed
## 7  Sally      3      C failed
## 8   Jeff      2      E failed
## 9   Jeff      4      C failed
## 10 Karen      3      C failed
```

4. Dates and Times with lubridate

```
library(lubridate)
```

```
## Warning: package 'lubridate' was built under R version 3.6.3
```

```
##
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
##
##      date
```

```
this_day <- today()
this_day
```

```
## [1] "2020-05-15"
```

```
month(this_day)
```

```
## [1] 5
```

```
wday(this_day)
```

```
## [1] 6
```

```
wday(this_day, label = TRUE)
```

```
## [1] Fri
```

```
## Levels: Sun < Mon < Tue < Wed < Thu < Fri < Sat
```

```
this_moment <- now()  
this_moment
```

```
## [1] "2020-05-15 17:59:16 IST"
```

```
second(this_moment)
```

```
## [1] 16.8503
```

```
my_date <- ymd("1989-05-17")  
my_date
```

```
## [1] "1989-05-17"
```

```
class(my_date)
```

```
## [1] "Date"
```

```
ymd("1989 May 17")
```

```
## [1] "1989-05-17"
```

```
mdy("March 12, 1975")
```

```
## [1] "1975-03-12"
```

```
dmy(25081985)
```

```
## [1] "1985-08-25"
```

```
ymd("1920/1/2")
```

```
## [1] "1920-01-02"
```



```
dt1 <- "2014-08-23 17:23:02"  
ymd_hms(dt1)
```

```
## [1] "2014-08-23 17:23:02 UTC"
```

```
hms("03:22:14")
```

```
## [1] "3H 22M 14S"
```

```
dt2<-c("2014-05-14", "2014-09-22", "2014-07-11")  
ymd(dt2)
```

```
## [1] "2014-05-14" "2014-09-22" "2014-07-11"
```

```
update(this_moment, hours = 8, minutes = 34, seconds = 55)
```

```
## [1] "2020-05-15 08:34:55 IST"
```

```
this_moment
```

```
## [1] "2020-05-15 17:59:16 IST"
```

```
this_moment <- update(this_moment, hours = 10, minutes = 16, seconds = 0)  
nyc <- now("America/New_York")  
depart <- nyc + days(2)  
depart <- update(depart, hours = 17, minutes = 34)  
arrive <- depart + hours(15) + minutes(50)  
arrive <- with_tz(arrive, "Asia/Hong_Kong")  
last_time <- mdy("June 17, 2008", tz = "Singapore")  
how_long <- interval(last_time, arrive)  
as.period(how_long)
```

```
## [1] "11y 11m 1d 21H 24M 16S"
```